Comments slides for Tuesday, November 17: Privacy and ethics, part 1

Matthew J. Salganik

COS 597E/SOC 555 Limits to prediction
Fall 2020, Princeton University

Moral dimension:

- Sociologists run away (or toward) from moral dimension

Moral dimension:

- ▶ Sociologists run away (or toward) from moral dimension
- ▶ Want research to be important, and moral basis is one way to do that

Moral dimension:

- ▶ Sociologists run away (or toward) from moral dimension
- ▶ Want research to be important, and moral basis is one way to do that
- ▶ Moral dimension in problem selection, not research design or data analysis
  (Arvind's point about communicating the findings is a good one for sociologists)

Moral dimension:
- ▶ Sociologists run away (or toward) from moral dimension
- ▶ Want research to be important, and moral basis is one way to do that
- ▶ Moral dimension in problem selection, not research design or data analysis (Arvind's point about communicating the findings is a good one for sociologists)
- ▶ Moral vs adversarial

Moral dimension:

- ▶ Sociologists run away (or toward) from moral dimension
- ▶ Want research to be important, and moral basis is one way to do that
- ▶ Moral dimension in problem selection, not research design or data analysis (Arvind's point about communicating the findings is a good one for sociologists)
- ▶ Moral vs adversarial
- ▶ Find something that works for you

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
    - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$

Kosinki et al.

▶ How to reconcile with the results of the Fragile Families Challenge?

  ▶ digital trace data vs survey data
  ▶ more cases
  ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
  ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
  - ▶ digital trace data vs survey data
  - ▶ more cases
  - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
  - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
    - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
    - ▶ r vs R^2. $r = 0.4 \to R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)
- ▶ Time stable vs time unstable traits

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
  - ▶ digital trace data vs survey data
  - ▶ more cases
  - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
  - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)
- ▶ Time stable vs time unstable traits
- ▶ Does test-retest reliability for say openness create a clear, sharp limit to prediction?

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
    - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)
- ▶ Time stable vs time unstable traits
- ▶ Does test-retest reliability for say openness create a clear, sharp limit to prediction?
- ▶ Unanticipated secondary use.

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
    - ▶ digital trace data vs survey data
    - ▶ more cases
    - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
    - ▶ r vs R^2. $r = 0.4 \rightarrow R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)
- ▶ Time stable vs time unstable traits
- ▶ Does test-retest reliability for say openness create a clear, sharp limit to prediction?
- ▶ Unanticipated secondary use. Are you worried about that for Fragile Families Challenge?

Kosinki et al.

- ▶ How to reconcile with the results of the Fragile Families Challenge?
  - ▶ digital trace data vs survey data
  - ▶ more cases
  - ▶ "prediction in the statistical sense" vs forecasting vs whatever we did in the Fragile Families Challenge
  - ▶ r vs R^2. $r = 0.4 \to R^2 = 0.16$ Why use r anyway?
- ▶ Learning model is simple (linear/logistic regression), data is key
- ▶ Lots of sparkle (e.g, likely curly fries predicts Raven's Standard Progressive Matrices)
- ▶ Time stable vs time unstable traits
- ▶ Does test-retest reliability for say openness create a clear, sharp limit to prediction?
- ▶ Unanticipated secondary use. Are you worried about that for Fragile Families Challenge? How does Kosinki's other work influence how you think about the ethics of this work?

Duhigg and Christl et al.

- ▶ Great lead

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?
- ▶ Target story part of the hype about machine learning.

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?
- ▶ Target story part of the hype about machine learning.
- ▶ Trace data to predict other trace data in the future (compare to trace data to predict survey data) (but why are birth records public?) (is this like nowcasting?)

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?
- ▶ Target story part of the hype about machine learning.
- ▶ Trace data to predict other trace data in the future (compare to trace data to predict survey data) (but why are birth records public?) (is this like nowcasting?)
- ▶ What if you are doing this stuff to get people to subscribe to the *New York Times*?

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?
- ▶ Target story part of the hype about machine learning.
- ▶ Trace data to predict other trace data in the future (compare to trace data to predict survey data) (but why are birth records public?) (is this like nowcasting?)
- ▶ What if you are doing this stuff to get people to subscribe to the *New York Times*?
- ▶ What if this was just measuring heterogeneity of treatment effects without inferring sensitive traits?

Duhigg and Christl et al.

- ▶ Great lead
- ▶ This seems creepy, but is it unethical? If so, why?
- ▶ Lack of transparency is a sign of a problem (adding in coupons for lawn mowers and wine glasses to confuse customer) "And we found out that as long as a pregnant woman thinks she hasn't been spied on, she'll use the coupons. She just assumes that everyone else on her block got the same mailer for diapers and cribs. As long as we don't spook her, it works." (Arvind example of FB sharing more information with advertisers than you)
- ▶ Arvind talked about privacy as a means to power/fairness/justice. Helpful for me. Is this a good strategy?
- ▶ Target story part of the hype about machine learning.
- ▶ Trace data to predict other trace data in the future (compare to trace data to predict survey data) (but why are birth records public?) (is this like nowcasting?)
- ▶ What if you are doing this stuff to get people to subscribe to the *New York Times*?
- ▶ What if this was just measuring heterogeneity of treatment effects without inferring sensitive traits?
- ▶ Can you measure not just tracking but harm?

Comments slides for Tuesday, November 17: Privacy and ethics, part 1

Matthew J. Salganik

COS 597E/SOC 555 Limits to prediction
Fall 2020, Princeton University